

# Accountant pattern: Lightweight solution for embedded micropayments

Jaro Šatkevič, me@JaroCoin.com

July 6, 2019

The growing number of decentralised platforms which require embedded payment mechanisms means that the need for a high throughput micropayments solution using cryptocurrencies is larger than ever before.

Presently, blockchain payments must be verified and stored by each node in the network. This means that sets the limit for the overall throughput of the system.

There are several off-chain protocols - *Lightning Network*, *Raiden* or various *Plasma* implementations. While they each provide solutions for specific use cases of micropayments with cryptocurrencies, they come with varying sets of limitations when it comes to high throughput micropayments in decentralized systems.

In this paper, we provide a construct of an embeddable, lightweight micropayment solution - "*Accountant pattern*".

Mysterium Accountant combines techniques used by payment hubs and digital cheque based unidirectional channels. This means the ability to provide up to couple of millions transactions per second. This best serves the needs of decentralized VPN, CDN or video streaming platforms.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Platform review and requirements for payment solution</b>	<b>4</b>
2.1	Components of the Network . . . . .	5
2.2	Initially proposed payment solution . . . . .	6
2.2.1	The fundamental issues with digital cheques . . . . .	7
2.3	Requirements for an “ideal” payments solution . . . . .	7
<b>3</b>	<b>Overview of potential solutions</b>	<b>11</b>
3.1	Blockchain per project . . . . .	12
3.2	Plasma . . . . .	14
3.3	Payment and state channels based solutions . . . . .	15
3.3.1	State channels . . . . .	16
3.3.2	Downsides and benefits of channels . . . . .	17
3.3.3	Micropayments networks . . . . .	18
3.3.4	Hashed Timelocked Contracts and Virtual Channels . . . . .	19
3.3.5	Known problems of micropayment networks . . . . .	22
3.4	Comparison and conclusion . . . . .	23
<b>4</b>	<b>Accountant pattern - lightweight payment protocol</b>	<b>24</b>
4.1	Payments over ”Accountant” . . . . .	24
4.2	Payment Promises based uni-directional channels . . . . .	26
4.3	Identity registry . . . . .	28
4.4	Off-chain messaging and promise exchange . . . . .	29
4.5	Incoming channel funds guarantees and trustless rebalance . . . . .	32
4.6	Transaction maker . . . . .	33
4.7	Properties of the protocol . . . . .	35
<b>5</b>	<b>Potential future work</b>	<b>35</b>
5.1	Migration into Connexor or Perun . . . . .	36
5.2	Plasma plus channels . . . . .	36
5.3	Take use of Counterfactual state channels . . . . .	37
	<b>References</b>	<b>38</b>

# 1 Introduction

The idea of building decentralized systems is not new. The main addition that cryptocurrencies have brought to decentralized protocols is the ability to enable incentivisation. This is done via embedded payment mechanisms.

In decentralized systems, it is often the case that two parties that do not trust each other want to carry out mutual transactions. This is especially relevant in the case of a decentralized VPN network.

## **Trustless transactions - the heart of a decentralized VPN network**

Considering the nature of a decentralized VPN network, one can assume that network participants would prefer to remain anonymous. Given the lack of trust, a service consumer will not be willing to pay a bigger amount up front. On the other end, the service provider isn't willing to deliver service without prepayment.

From this above mentioned scenario arises a need to split the service into chunks - provided in exchange for micropayments - meaning that each party risks a much smaller value. Put into context within a decentralized VPN network, this would mean that a user will pay a node a couple of times a minute, sending a tiny amount of tokens in exchange for the bandwidth they are renting.

With cryptocurrencies, it is now possible to embed decentralized payment solutions in a trustless and permissionless way. Unfortunately, because of technical limitations, to make scalable cryptocurrency payments, additional work is required.

The backbone of cryptocurrencies is a technology called the *blockchain*. This technology requires each transaction to be stored on the ledger and replicated across thousands of network participants. This imposes a fundamental limit on the number of transactions that can be processed.

Let's take two of the most popular blockchains, Bitcoin and Ethereum. Their blockchains can provide throughput of 3 to 25 transactions per second causing high transaction fees and making on-chain transactions too expensive for micropayments.

To overcome blockchain scalability problems, several approaches have been proposed:

- ***Speedup blockchain*** itself by changing some fundamental components.

- Use *blockchain attuned to use-cases* (sidechains, Cosmos and Polkadot networks).
- *Plasma* and other kinds of child-chain solutions.
- Payment and *state channels* based solutions.

Most of the proposed solutions are in extremely early stages of development and adoption. Each come with their own pros and cons and there is currently no clear industry standard.

This lack of payment standardization is one of the contributing reasons to decentralized protocols still being early betas. While some may already have working core solutions, payments are often a missing and evolving piece.

**The goal of this whitepaper is to lay out research on the simplest, cheapest and powerful solution for high throughput micropayments using tokens.** This is in the context of creating a payments infrastructure for decentralized VPN, CDN or streaming services.

Here is the methodology we have undertaken to understand the problem and lay out the solution for high throughput micropayments:

1. Analyse the needs of decentralized VPN, CDN or streaming platforms, collecting requirements for payment solutions which are applicable.
2. Research and compare advantages, limitations and side effects of current layer two solutions. We will look at state channels, micropayment channel networks, plasma and sidechains. In this research, we will dive into the possibility of the use of each of these layer two solutions as a base for high throughput micropayments in decentralized platforms.
3. Detail a lightweight and easily embeddable solution for high throughput micropayments.

## 2 Platform review and requirements for payment solution

Let us take the case of Mysterium Network [5] - the world's first decentralized VPN. This is a decentralized application which requires high throughput micropayments.

Mysterium Network is a network of nodes providing security and privacy for end users of MysteriumVPN, a decentralized VPN application.

The goal of Mysterium Network is to combine powerful encryption mechanisms, reputation management systems and layered protection protocols to build an infinitely scalable P2P architecture. A decentralized VPN is fundamentally different from current in-market VPN solutions in that its architecture is fundamentally decentralized. The long term vision of Mysterium Network is to build a system where there is no single point of failure in the network. This is possible as the network, in the case of Mysterium Network, is a community of incentivised nodes renting their bandwidth in exchange for payment.

This means that no government, company or person will be able to censor or stop the service. Even the team of Mysterium Network will be not able to censor or stop the network once the final solution has been released into production.

## 2.1 Components of the Network

- **Provider nodes:** This is a service written in golang which can be run by any user of the network. This service allows you to earn cryptocurrency for renting your bandwidth to other users (consumers) using OpenVPN.
- **Consumer app:** This is a dApp written in golang and JavaScript which can be run by any user of the network to connect to provider nodes via OpenVPN and rent their bandwidth.
- **Discovery service:** Discovery service helps providers to advertise themselves in the network. This allows consumers to find them and rent their bandwidth.
- **Myst token:** This is the utility token used within Mysterium Network. MYST is the main and only method of P2P payment within the network. MYST token is released as an ERC20 token on Ethereum blockchain.
- **Mysterium ID:** Each user of the network (consumers and providers) has to register a Mysterium Identity, which is similar to an Ethereum address (20 bytes of keccak256 hash of the public key derived from private key using elliptic curve cryptography).

Mysterium ID is needed to:

- Uniquely identify network actors;
- Be able to securely establish connections between them;

Mysterium ID has to be registered on-chain using smart contracts deployed into Ethereum network.

## 2.2 Initially proposed payment solution

Initially, the Mysterium team was planning to use a mechanism which remotely resembled the way cheques work. A blockchain account holder can write a cryptographic cheque to another account (the beneficiary) as a form of payment. A cheque would include the issuer's address, beneficiary's address, the sum of the total promised, the sequence and signatures. The amount is written on the cheque can be updated and only the last version of the cheque is valid. The beneficiary can settle the promised account on blockchain at a later date.

Here is a snippet from the smart contract code which settles the promised value on-chain.

```
1  function settlePromise(address issuer,
2                          address beneficiary,
3                          uint256 seq,
4                          uint256 amount,
5                          bytes issuerSignature,
6                          bytes beneficiarySignature) public
7  {
8      bytes32 promiseHash = keccak256(beneficiary, seq,
9                                       amount);
10
11      address recoveredIssuer = ecrecover(promiseHash,
12                                          issuerSignature);
13      require(recoveredIssuer == issuer);
14
15      address recoveredBeneficiary = ecrecover(promiseHash,
16                                              beneficiarySignature);
17      require(recoveredBeneficiary == beneficiary);
18
19      require(seq > clearedPromises[issuer][beneficiary]);
20      clearedPromises[sender][receiver] = seq;
21
22      require(token.balanceOf(issuer) >= amount);
23      token.transferFrom(issuer, beneficiary, amount);
24
25      emit PromiseSettled(issuer, beneficiary, seq, amount)
26      ;
27 }
```

This solution is much better than doing on-chain transactions for each layer of this transaction.

Cheques with promised amounts can be sent from consumer to provider peer to peer each second. Only the *consumer* and *provider* will be aware of these changes in states, until the *provider* finally decides to settle transactions on the blockchain. This will drastically reduce the amount of on-chain transactions while providing a secure way for two parties that don't trust each other to transact.

### 2.2.1 The fundamental issues with digital cheques

In this section we will define what a digital cheque is, and explore fundamental flaws that have yet to be fixed.

A digital cheque is a cryptographic signature which can be sent into smart contract to prove ownership of tokens and empowers smart contract to send tokens into given in that signature address.

Here are some of the issues we face with digital cheques:

#### 1. Double spending

When there isn't enough funds in the issue's balance to cover the promised value, the transaction will be rejected when it is settled into the blockchain. This creates the possibility of double spending. A consumer may be able to issue a cheque for the same funds to another party.

In the case of a VPN service, this isn't a huge problem as bandwidth is a "*perishable product*". However, this risk of double spending forces more regular on-chain settlements. Service providers would be settling each time the promised amount was big enough that the risk of losing that value outweighed the cost of the transaction.

#### 2. Small value cheques

Within the decentralized VPN system, a consumer could use a provider's service only once and for a limited time (e.g. one hour unblocking Netflix). This means that providers will accumulate a lot of small value digital cheques (e.g. 0.10 USD value in tokens). Settling one cheque on the Ethereum blockchain can cost from 0.01 up to 1 USD depending on network congestion. This means that either large fees will be paid or these cheques will remain unsettled.

## 2.3 Requirements for an "ideal" payments solution

Due to the high level of privacy and anonymity in decentralized networks, actors within the network do not trust each other. Additionally, there is

no trusted intermediary available which can act as a custodian or help with conflict resolution. In these situations, the consumer isn't going to pay a large amount up-front and the service provider will be unlikely to offer services without prepayment.

In the case of a decentralized VPN, a service can be split into microservices and provided in exchange for micropayments, reducing the risk for both parties.

This *pay-as-you-go* model means that both parties begin transacting immediately, with transactions occurring a couple of times a minute while sending small amounts of cryptocurrency.

### 1. Consumer to Provider payments

Usually there are two actors in the network, consumers and providers of a service. All payments are sent by consumers and received by providers. Never vice versa.

Thanks to this observation, the protocol can be simplified and make use of *payment promises* or *uni-directional* micropayments channels.

### 2. High throughput and scalability

It is extremely important that the protocol is able to support frequent small payments (e.g. 10 seconds by all participating parties of values less than 1 cent). This means that:

- The payment solution should be able to process as many transactions per second as there are active sessions established between service providers and consumers.
- Value of a transaction should be set to parts of a cent (in given cryptocurrency).
- Transactions should be marked as final in a very short period of time. They ideally should have an *instant finality* property.
- Fast answers about transaction status with minimal networking errors and retry amounts.
- Transaction fees have to be very small (parts of a penny) or expressed as a percentage of the transaction value.
- There should be minimal presence on-chain. This system would need to aggregate payments of many sessions and from different consumers, to settle them in one transaction on chain.

### 3. Utility token and stable coin support

Bitcoin, Ether or other popular cryptocurrencies are volatile and may



not be accepted as payment in certain cases. Further there are several decentralized apps which have issued their own utility token (*MYST* in the case of Mysterium Network). This means there is a requirement for payment protocols to support transactions using *ERC20* tokens issued on Ethereum blockchain.

#### 4. Security

Digital services such as VPN can be seed as *perishable* as they are renting traffic which if not used, is “gone” forever. This means that some level of risk of unpaid service can be acceptable. However, it should be up to the service provider to weight this level of risk against convenience (usability and performance) and cost (lower on-chain settling fees). In a general case, this double spending attempt has to be immediately identified and such transactions, rejected.

Anonymity, and the permissionless nature of decentralized systems create cases where additional layers of protection against bad actors is required. For example, in the case of a decentralized VPN, poorly performing service providers may simply refresh their identities by simply creating a new identity in the network. Bad acting consumers could organise DDoS attacks, and it would be impossible for providers to ban them as they too could continuously create new identities.

To curtail this behaviour we need a reputation system which incorporates identity registration, staking and a punishment system.

- network identity have to be registered in given smart-contract, to do so there should be paid registration fee or staked given amount of tokens;
- there should be not possible to pay with same coin twice (avoid double spending);
- system should be secure against different kind of attacks (e.g. DDos);

#### 5. Decentralization

Components as important as payments should have a high degree of liveness.

There should be no central party which is easy to shut down and censor. This means that the payment protocol that is needed should maintain at least some level of decentralization and should not be operated by any party.

## 6. Low complexity implementation

There are several potential ways we can embed payment solutions into decentralized platforms. We could either progress with a popular and scalable payment network which already works on several platforms, has stable APIs and is already in use in many communities.

Or we could work with a solution with a low level of complexity so it could be easy to implement, cheap to operate and resource efficient to use.

- (a) **Easy implementation.** Such protocol needs to be used and modified to suit the needs of various systems. This means it is necessary that each of it's main parts can be reimplemented in main programming languages in a matter of weeks.
- (b) **Cheap to operate.** To make solutions easier to decentralize, it should be relatively easy and cheap to run and operate nodes. No big initial financial investment, complicated installation and advanced hardware should be used.
- (c) **Efficient usage.** Since there is a requirement to create a stable and fast payment solution, the less communication messages that need to be exchanged between parties, and the less intermediaries involved, the better for the protocol.

## 7. Good user experience

Technical parameters alone do not make for successful systems. If there is friction involved in a user making a payment, we will not be geared towards adoption.

Here are various usability considerations that we need to solve for:

- (a) Consumers should be able to deposit funds using any popular crypto wallet, or directly from exchanges.
- (b) Users should have the possibility to own just one asset (e.g. MYST token) and not be required to own any additional utility token or coin to make payments within the network. This is especially important when considering the ERC20 nature of MYST, and the need to settle on the Ethereum blockchain, which requires a fee in ETH for each transaction.
- (c) Any cryptographic proof should be able to be transferred not only by a signing party, but by any third party. All while maintaining the same level of security. **This allows the use of cloud trustless services to send, valid and protect transactions.**

- (d) Since service providers are earning income within the system, they can be required to stake tokens or pay a platform fee. Consumers however need the freedom to use the network without staking.

### 3 Overview of potential solutions

*Blockchain* is a replicated state machine which orders transactions on it's global state. Transactions are verified and replayed by each participant, called a *node*, in the network. This limits the throughput of the network as a whole to the lowest throughput of any of it's nodes. Increasing the load beyond that throughput may result in nodes unable to handle the load being pushed out of the network. This impose a fundamental limit of amount of transactions which can be processed.

Let's take two most popular cryptocurrencies Bitcoin and Ethereum. Their blockchains can provide throughput of 3 to 15 transactions per second which in busy period causes high transaction fees and makes on-chain transactions too expensive for micropayments.

One of solutions would be to increase block size which could allow increase throughput. There are blockchains with block limits of 128Mb (100 times higher than bitcoin's) so theoretically they can process up to couple of hundreds transactions per second. Such level of scalability work for one-time type of payments (like buying cup of coffee) but for high throughput nano payments we need solution which could process millions of transactions per second. Unfortunately there are physical network, validation time (CPU) and disk space limits so blocks can't be ultimately big.

Finally there is one more problem. Participants in the network by default agree that the chain with the highest *difficulty* or more blocks put into it, is "true ledger". If some other branch mined in parallel will get longer chain, a new ledger is accepted as the true ledger. Some transactions accepted into older ledger can be not added into new one. This means that the longer waiting period before accepting payment, the bigger guarantee the transaction written in blockchain is irreversible.

For frequent transactions, the faster *finality* there is, the better. *On-chain* transactions however don't have reasonably fast finality.

To overcome *on-chain* limitations, such as scalability, high fees and long finality time, several *off-chain* protocols (or so-called *Layer 2* solutions) have been proposed. Among these techniques, one important differentiator is whether the relocated operations introduce additional consensus assumptions (like in sidechains and interoperable blockchain networks such as Polkadot

or Cosmos), or allow users to restore their state to the original blockchain (like in state channels or Plasma).

### 3.1 Blockchain per project

The most radical solution would be to move platform specific utility token into own blockchain. There are couple of frameworks to build own blockchains with pluggable consensus and application layer. Most promising solutions are *Tendermint* developed by Cosmos Network, *Substrate* developed by Parity Technologies as part of Polkadot Network, using software of Ethereum but with different consensus algorithm and *Hyperledger Fabric* by Linux Foundation.

Nevertheless that these solutions are relatively hard to customise and launch, using them would additionally require building community of full nodes and stake pools which would guaranty network's security. Another unwanted side effect would be need of moving token into new blockchain. This is socially complicated and hard to implements project.

Alternatively there could be launched a *sidechain* - a blockchain ledger that runs in parallel to a primary blockchain. Assets from the main blockchain can be linked to and from the sidechain. This allows the sidechain to operate independently of the primary blockchain by introducing own consensus mechanism, faster speed of transactions and features needed to some dedicated platform.




Figure 1: Two-way pegged sidechain

Thanks to possibility to use own, less decentralised but more scalable consensus algorithms, *sidechains* can provide much higher throughput than

in primary blockchain. However while introducing consensus assumptions which in situation of fail will permanently compromise long-term guarantees (such as persistence of asset ownership). If state is "moved" to a sidechain and that chain's consensus mechanism fails, owners or beneficiaries of that state may lose everything delegated there, even when the primary blockchain remains secure.

Slightly different approach is taken by interoperable multi-chain networks such as *Polkadot* (see figure 2) which allows new designs of blockchains (called parachains) to communicate and pool their security while still allowing them to have the entirely arbitrary state-transition functions. This helps to bootstrap new chain much faster while having same security guaranties as whole Polkadot network.




Figure 2: Polkadot network components

Despite the fact that *sidechains* can provide a significant increase of transaction throughput and in case of *multi-chains* also relatively easy achievable high level of security guarantees, general issues of blockchains still apply. Because of blockchain nature and physical networking and storage limitations, they still can't provide throughput of hundreds of thousands transactions per second and *instant finality*, which may be needed for *micropayments* in successful decentralised VPN or streaming platform.

## 3.2 Plasma

*Plasma* [8] is a proposed framework for scaling Ethereum capacity by using hierarchical sidechains. Plasma type of sidechains (also called child chains) allow to do a majority of transactions outside of the "root chain" (e.g. Ethereum). Only deposits and withdrawals, the points of entry and exit, are handled on the root chain smart contract.

Similar to blockchains Plasma chain stores all its transactions packed into blocks while using *UTXO* for balance accounting. To make sure that transactions are final, Plasma operator is doing something called a "state commitment". It is a cryptographic way to store a compressed version of the state of child-chain inside of root chain. Typically all the state is stored in *merkle trees* and only *merkle root* of each block's state is persisted into root chain.

Usually Plasma chains are run by single operators, having distributed nodes using some kind of Byzantine Fault Tolerant (BFT) consensus (e.g. Proof of Stake) is possible though.

Differently to sidechains and multi-chains, Plasma makes possible for users to leave at any time. Such action is usually referred as "exiting". This allows users safely withdrawal their funds out of Plasma even if it was shut-down by operator.

While it offers significant speed (up to 1000 tx/s) and latency improvements over Ethereum itself, Plasma cannot offer the near-zero latency and near-free transaction fees. It also requires significant storage to store its ledger (especially with big amounts of transactions).

Another downside of Plasma chain is that it is complex and difficult to implement solution. Especially hard is to run it with distributed nodes which have BFT type of consensus. When Plasma chain is operated with single operator, there is risk that he will create 'fake' blocks, which may end up with mass exits out of plasma into root chain and stuck of whole network.

One more downside is that plasma is quite expensive to operate. Each couple of blocks it have to commit its state to root chain, which means, that plasma operator would have to pay 1080 USD (as gas fee) per day.

$$3 \text{ tx/minute} * 1440 \text{ minutes/day} = 4320 \text{ tx/day}$$

$$1 \text{ tx} = 25 \text{ cents}$$

$$4320 \text{ tx/day} * 0.25 \text{ USD/tx} = 1080 \text{ USD/day}$$

This isn't big problem when there is already significant load in the network, however in the beginning covering such costs can be problematic.

And finally, regardless of that Plasma transaction throughput is significantly higher than Ethereum's, it is still not enough for distributed VPN network needs. As alternative, there could be possible to build some kind of payments channels on top of Plasma (see 3). In this way parties could do peer-to-peer transactions while having active service and close channels right after closing connection.




Figure 3: Plasma + payment channels

In situation when there are more than one transaction per second throughput, Plasma transactions can be relatively cheap (less than 1 cent per transaction) and can have fast including into block (from 1 to 10 seconds, depending on implementation). This allows opening and closing channels when needed and avoiding long waiting times.

Unfortunately this solution is even more complicated than using just plasma and has similar cost and decentralisation issues mentioned above.

### 3.3 Payment and state channels based solutions

A *micropayment channel* is a class of techniques designed to allow parties to exchange digital value without committing all of the transactions to the blockchain. In a typical payment channel, only two transactions are added to the blockchain but an unlimited or nearly unlimited number of payments can be made between participants.




Figure 4: Sequense based payment channel

To open payment channel parties have to lock some funds into multisig-nature smart contract (see figure 4). This allows parties to updated channels balances off-chains while beeing sure with hight probability that funds will be not double spent or stolen.

There are various payment channels techniques such as sequence based channels, duplex channels, time-locked channels, etc. Any of them would improve security and would help to avoid double spending problems comparing to initially proposed digital cheques solution (see section 2.2). Also that would allow providers to wait longer until settling on-chain what could reduce settling costs because some consumers may return to use service via same provider again.

### 3.3.1 State channels

State channels are the general form of *payment channels*, applying the same idea to any kind of state-altering operation normally performed on a blockchain. Moving these interactions off of the chain without requiring any additional trust can lead to significant improvements in cost and speed. State channels will be a critical part of scaling blockchain technologies to support higher levels of use.



The basic components of a state channel are very similar to that in payment channels.

1. Part of the blockchain state is locked via multisignature type of smart contract, so that a specific set of participants must completely agree with each other to update it.
2. Participants update the state among themselves by constructing and signing transactions that could be submitted to the blockchain, but instead are merely held onto for now. Each new update "trumps" previous updates.
3. Finally, participants submit the state back to the blockchain, which closes the state channel and unlocks the state again (usually in a different configuration than it started with).

Because tokens on Ethereum blockchain are represented in form of state in smart contracts, state channels are the way how to move token interaction into *layer 2*.

### 3.3.2 Downsides and benefits of channels

The downside of using payment channels is that parties have to lock funds into multisignature smart contracts which in comparison to *digital cheques* means additional on-chain transaction for opening channel and needs of having locked bigger amounts of funds. There is not possible to reuse same funds in another channel before closing previous one. Another downside is time to service. Opening new channel will take some time (depending on type of blockchain and network load this make take from one minute till up to couple of hours) which makes bad user experience.

There are also couple of additional benefits of payment channels:

- **Privacy** – each transaction is known only to participating parties, which allows them to have privacy in their individual transactions such that the public only knows the result which has been shared on-chain.
- **Instant finality** – parties can sign and exchange messages instantaneously without having to wait for confirmation from the blockchain. This is leading to a much better user experience.
- **Lower cost** – digital value is exchanged with the only a few on-chain transactions being made for creating the channel and for closing the channel.

### 3.3.3 Micropayments networks

Payments channels is very promising technique for micropayments. However opening new channel requires on-chain transaction which is both slow and expensive and can't be done in setups when there are a lot of service providers. In this situation payment channels based solution is reasonable only when participants do not need to be connected to everyone else (see figure 5).




Figure 5: Everyone makes channel with anyone else

Many existing research on payment channels have focused on exploring the design space of how to structure payments through intermediaries [2, 3, 7]. Let's suppose that Alice has a payment channel with Ingrid, and Ingrid has one with Bob. If Alice pays Ingrid off-chain and then Ingrid pays Bob the same amount, this is equivalent to Alice paying to Bob off-chain, without requiring a new Alice-Bob payment channel to be set up.

In situation when parties don't have channels with single intermediary, a micropayment channel network can be used together with a routing algorithm to send funds between any two parties in the network (see figure 6).




Figure 6: Micropayment channels network

As micropayment channel network can keep most transactions off-chain, blockchain based currencies may scale to magnitudes larger user and transaction volumes than any currently existing centralised solution. Also, micropayment channel networks allow for fast transactions thanks to *instant finality* property of payment channels. Transaction is final as soon as it is signed and sent to another party, the blockchain latency does not matter.

### 3.3.4 Hashed Timelocked Contracts and Virtual Channels

There are various ways to create micropayments networks trustlessly (i.e. to ensure that Alice pays Ingrid if Ingrid pays Bob). The most popular is Hashed Timelock Contract (HTLC) based approach [7]. In its canonical way an equal amount of funds from both payment channels are locked up in a way that they can only be spent if a certain hash is revealed before a certain deadline (thus, locked “by hash” and “by time”).

This technique can allow payments to be securely routed across multiple payment channels and is used in Bitcoin’s Lightning Network and Ethereum’s Raiden Network.




Figure 7: HTLC based payment - Alice pays Dave

In figure 7 we see atomic value exchange over three micropayment channels (Alice to Bob, Bob to Carol, Carol to Dave). If Alice will pay to Dave via Bob and Carol, she needs to ensure that Bob and Carol cannot run away with her money. To do so, Dave generates pre-image  $R$ , which is random number (shown as key on scheme), and then shared with Alice hash  $H$  (shown as lock on scheme) of this pre-image. In step 2 Alice generates an HTLC with Bob that says *"I will pay you  $X$  coins if you show me the pre-image  $R$ . If you don't show  $R$  during period  $\Delta t$ , I will take back my coins."*. In step 3 Bob generates similar HTLC with Carol, but sets period of  $\Delta t - 1$ . In 4 step Carol do same with Dave, but with even shorter period of  $\Delta t - 2$ . Then in step 5 Dave will show  $R$  to Carol, in next step Carol shows  $R$  to Bob and finally Bob shows it to Alice. At this point payment of  $X$  amount from Alice to Dave is finalised.

HTLCs can be established in any chain of any length consisting of different payment channels. As an incentive for intermediate hops to forward transactions a small fee can be charged for using the service of the channel. Fee payments are also justified as the balance on a channel gets shifted, which is only beneficial for balancing a lopsided channel. After a successful transaction with HTLCs, channel parties do not need to broadcast their contract and can just replace their HTLC with a new commitment transaction without an HTLC. HTLCs can be combined with timelocks or revocable transactions changing the output of the HTLC accordingly.

In Ethereum blockchain, which supports advanced smart contracts, routing payments could be done in an alternative form via a technique known as *Virtual channels*, introduced by *Perun* paper [3] and simultaneously similar technique was proposed by *Counterfactual* protocol (named as Meta chan-

nels), which use an intermediary that serves as a *"virtual payments hub"*. Anyone with a payment channel connected to the hub could establish virtual channels between each other (see figure 8). Unlike routing payments via HTLCs, Hub does not need to be involved in every payment between Alice and Bob. This property reduces latency and costs, while increases privacy.

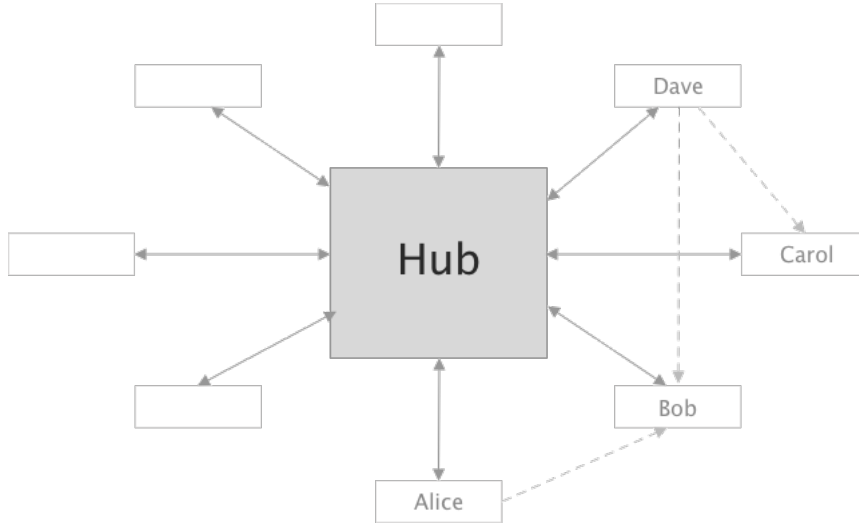


Figure 8: Virtual Payment Channel Hub

To open virtual channel, Alice and Bob essentially need to lock-up a set number of coins from their payment channel with Hub. The amount of locked coins will become the value of the virtual channel between Alice and Bob. Notice, that the Hub remains financially neutral by simply mirroring balances.

This technique could be repled to increase the length of the channel (see figure 9). Because a virtual channel is an instance of an off-chain contract, increasing the length not require another transaction on-chain.




Figure 9: Bob becomes the Hub for channel between Alice and Carol

Virtual channels also represent a different business model from channel routing. HTLCs have a “*pay-per-payment*” fee model where there is need to incentivize each intermediary to route each payment. Virtual channels, on the other hand, have a “*rent-a-path*” fee model. In this model, an intermediary acts as a virtual payment hub that has direct channels with multiple parties. If Ingrid is the intermediary, then Alice and Bob pay Ingrid to keep the channel open for a certain period of time. Such a model might have better economics for high-volume micropayments.

While Perun and Counterfactual have introduced novel ways of routing payments across multiple intermediaries, HTLC-based routing is currently the primary implementation on blockchain mainnets, is better tested and much easier to implement.

### 3.3.5 Known problems of micropayment networks

Micropayment channels networks looks really promising but unfortunately they have couple of negative sides.

Downside of using HTLC based micropayment networks is complicated routing. *Lightning Network* is great for time to time low value payments, however it is much harder to use for frequent transactions into same receiver because of need of constantly checking if all parties in selected route are still alive and good to be used (e.g. all channels has enough funds to send payment forward).

Complicated routing issue could be solved by using payment hub model, when one of a few intermediaries have a lot of channels with end users. Downside of this solution is that this requires the Hub to lock a lot of funds in channels with all actors in the network. Another downside is that this solution is centralised and if hub’s operator will be required to stop activities or will decide to shutdown servers by any reason, the whole network’s payments activities will be stopped.

Another critique of payment channels network is that both parties (paying and receiving) have to be on-line during payment. This issue however is not valid for high frequency nano payments in decentralised networks, because we’re sure, that payer (service consumer) and payee (service provider) are on-line, otherwise service wouldn’t be provided and there would be no need for payment.

Finally, micropayment channel networks are really useful when there are a lot of people already using them. And when there are huge support for all popular programming languages and platforms. Unfortunately, the most popular solution (*Lightning Network*) don’t support tokens. On Ethereum there are a few promising projects (e.g. Raiden, Connex, Perun or Coun-

terfactual), but they either are not widely adopted, either in early stage of development. Because they don't have rich tooling (e.g. protocol client implementations in popular languages like golang, rust or python) and are changing rapidly, they are quite complicated to implement and maintain. There are no clear winner at the moment and trying to stick into one of general purpose payment network's could mean choosing wrong direction which may mean either becoming main network developer, either having complicated migration into another (much more popular at that time) solution in the future.

### 3.4 Comparison and conclusion

The key insight of *Layer 2* solutions is that not every transaction has to be applied globally. All analysed solutions solves this problem in different ways and has different trade-offs.

Many state channel developers see *Layer 1* as the Security Layer and *Layer 2* as the Scalability Layer. Furthermore, Layer 2 provides lower latency and cost per transaction that would otherwise not be possible beyond a certain level of throughput with Layer 1 solutions.

Table 1: Comparison of researched solutions

	Channels	HTLC Network	Payment Hub	Plasma	Sidechain
On-chain transactions	●●●	●●	●●	●	●
Instant finality	✓	✓	✓		
High throughput	●●	●●●	●●●	●	●
Off-chain messaging	●	●●●	●●	●	●
Decentralised	✓	✓	±	?	±
Requires specialised wallet	?	✓	✓	✓	✓
Topup from exchange	?	—	—	—	✓
Easily embedable	✓	±	—	±	±
Overall user experience	●	●	●●●	●●	●●
Implementation complexity	●	●●●	●●	●●●	●●●
Fast settlement	?	—	±	✓	—
Security	●●●	●●●	●●●	●●●	?
Utility token support	✓	✓	✓	✓	✓

Explanation of meaning of used symbols: ●●● – high, ●● – medium, ● – low, ✓ – yes, ± – more or less, ? – unknown, depends on implementation.

Explanation of compared solutions:

- **Channels** – pure payment channels based solution (without network), when each pair of consumer and provider have to open new channel.

- **HTLC Network** – Micropayment channels networks which is using hashed timelocked contracts for payment routing, e.g. Raiden Networks.
- **Payment Hub** – Virtual channels hubs based payment solution, such as Perun, Counterfactual or Connex.
- **Plasma** – One of Plasma chain implementations (e.g. Plasma MVP), dedicated to serve for payment purposes for particular distributed system.
- **Sidechain** – Running sidechain using Substrate or Tendermint frameworks and connected to Polkadot or Cosmos.

In table 1 we can see comparison of properties given by analysed solutions. It is clear, that because of *instant finality* and higher throughput possibilities, micropayment channels based solutions (such as state channels, htlc networks or virtual payment hubs) fits better for high frequency micropayments in decentralised platforms.

## 4 Accountant pattern - lightweight payment protocol

After analysing currently available options became obvious that at this stage it may be better solution to introduce own, lightweight, state channels based, easy to implement and maintain solution.

In this section idea of "*Accountant pattern*" protocol will be described. Essentially it is solution taking best parts of Payment promises, uni-directional channels and paying using single intermediary taking use of hashed time-locked contracts for avoiding need of introducing trusted custodian.

### 4.1 Payments over "Accountant"

Digital cheques or *Payment promises* is elegant, efficient and easy to implement solution. Thanks to some modifications (comparing to initial solution, described in 2.2 section) they have to be signed by only one party, are easy to verify, don't require complicated setup procedure, have practically unlimited amount of updates and require to store only last version of issued promise.

Main problems of using them is possibility to do double spending which pushes for more often on-chain settlements. To solve this problem there is suggestion to introduce one more party called *Accountant* which should see



and verify promises issued by consumer and be aware of actual balance of consumer's funds (see figure 10).




Figure 10: Payments via Accountant

Accountant is the most complicated to implement part of this protocol. He have to store state of balance of each consumer and be able to accept many requests to validate if transaction is valid.

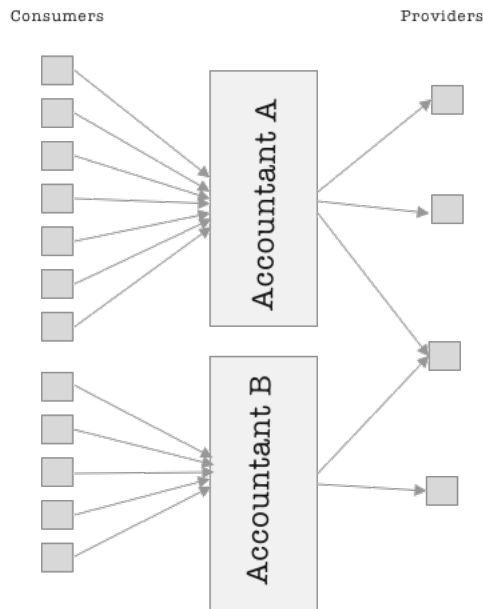


Figure 11: Multiple accountants may be used in network

To maintain requirement of decentralisation there may be deployed multiple accountants in the network (see figure 11). Each consumer usually works with only one accountant while providers may need to be aware of different accountants. Accountants are chosen by consumers but because of cryptographic schemes described in next sections they are non-custodian and can't steal any funds or cooperate with consumers to cheat against providers.

## 4.2 Payment Promises based uni-directional channels

To organise secure, non-custodian and trustless work via accountant, we have to maintain two types of channels: paying channels (consumer  $\rightarrow$  accountant) and receiving channels (accountant  $\rightarrow$  provider). So accountant plays similar role as intermediary or hub described in section 3.3.3.

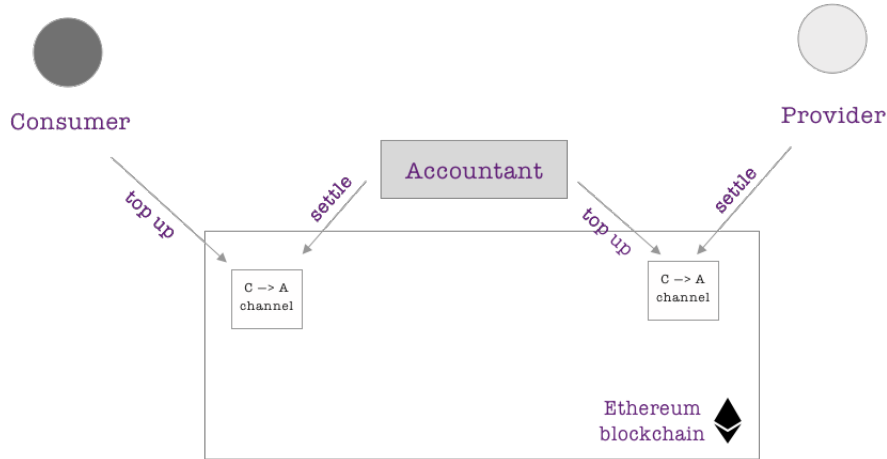


Figure 12: Two type of channels

Because in our case we have situation when payments are done in one direction (from consumer to provider) we may organise uni-directional state channels which are very similar to payment promises and have most of their properties. The only differences are that there should be done accounting and funds freezing for each channel separately. Also hashlock check (HTLC described in section 3.3.4) was added to guaranty that Accountant will not keep funds to itself.

Implementation of promise settlement function can be found above and full state channel smart contract can be found in github repository [6].

```

1  struct Party {
2      address beneficiary; // funds destination
3      uint256 settled;     // total amount already settled
4  }
5
6  function settlePromise(uint256 _amount, bytes32 _lock,
7      bytes memory _signature) public {
8
9      bytes32 _hashlock = keccak256(abi.encode(_lock));
10     address _channelId = address(this);
11
12     address _signer = keccak256(abi.encodePacked(
13         _channelId,
14         _amount,
15         _hashlock
16     )).recover(_signature);
17     require(_signer == operator);
18
19     // Calculate amount of tokens to be settled.
20     uint256 _unpaidAmount = _amount.sub(party.settled);
21     require(_unpaidAmount > 0);
22
23     // If signer has less tokens than asked to transfer,
24     // we can transfer as much as he has already and rest
25     // tokens can be transferred via same promise but in
26     // another tx when signer will topup channel balance.
27     uint _currentBalance = token.balanceOf(_channelId);
28     if (_unpaidAmount > _currentBalance) {
29         _unpaidAmount = _currentBalance;
30     }
31
32     // Increase already paid amount
33     party.settled = party.settled.add(_unpaidAmount);
34
35     // Send tokens
36     token.transfer(party.beneficiary, _unpaidAmount);
37 }

```

Because our channels are uni-directional instead of sequence number there is simply used total promised amount. Smart contract is calculating difference between amount on given promise and already settled amount.

$$amountToTransfer = totalPromised - alreadySettled$$

This gives high level of flexibility. Promises can be settled in any order with any gaps (e.g. only one of thousand promised can be send to settle)

and amounts still will be calculated properly.

### 4.3 Identity registry

Another important component of protocol is *registry* where all identities (both, consumers and providers) and accountants should be registered. During registration there will be also created payment channel between consumer and accountant (see figure 13).

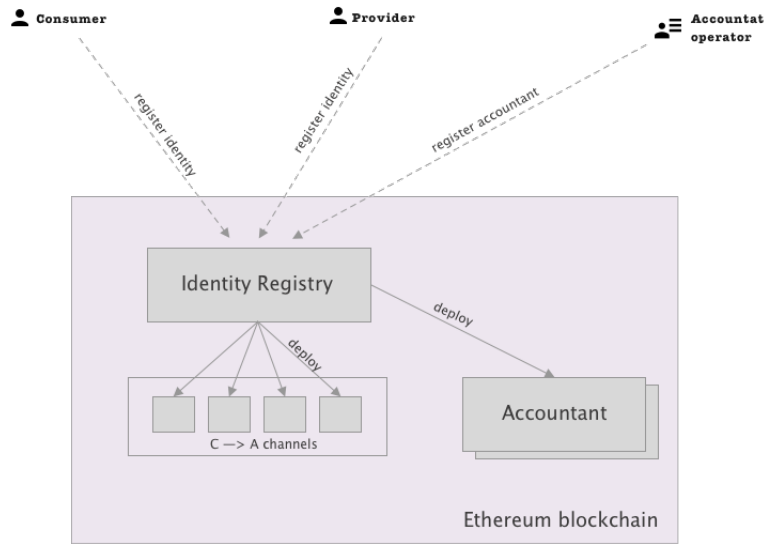


Figure 13: Registration of accountant and identities

Because there will be need to deploy channel per each registered identity there have to be done optimisation in terms to save of channel deployment transaction fee. This protocol suggests to use *EIP1167* (Minimal Proxy Contract [4]) which will proxy all requests into single *ChannelImplementation* contract while still allow to maintain separate state and address per each channel (see figure 14).

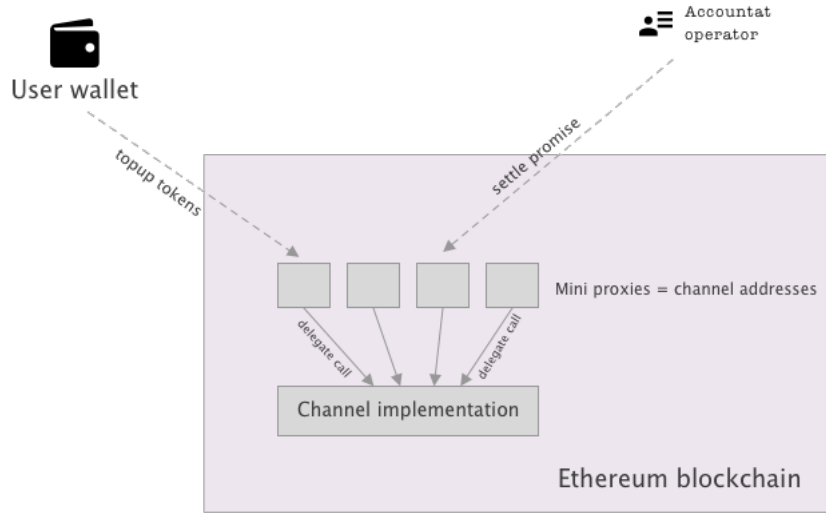


Figure 14: Minical proxies pointing into single implementation

Consumer (paying) channels are deployed using *CREATE2* opcode [1] which allows deterministically calculate future channel's address.

$$keccak256(0xff + registry + identity + keccak256(byteCode))[12 :]$$

This means, that users can topup their channels even before identity was registered and actual smart contract code was deployed there. Also because each channel has separate address and it's topup don't require adding any additional parameters (payload) to the transaction, they can be topuped using any wallet with token support or directly from exchange.

#### 4.4 Off-chain messaging and promise exchange

Differently than in *Lightning Network* or *Hub* based payment networks, in *Accountant pattern* payments are not going through intermediary, but are "verified" by Accountant. There is used payment promise exchange technique (see figure 15).




Figure 15: HTLC based payment - Bob pays Alice

Payment is finalised in five steps. At first Provider have to *generate invoice* and send it to Consumer.

$$invoice = [hashlock, agreementId, agreementAmount]$$

$$hashlock = keccak256(R)$$

$$R = HugeRandomNumber$$

Then Consumer issues payment promise, which allows Accountant to settle given amount of funds, pack it into so called *Exchange message* and send back to Provider.

$$Message = keccak256(channelId, totalPromisedAmount, hashlock)$$

$$Signature = sign(Message)$$

$$Promise = [Message, Signature]$$

$$ExchangeMessage = [Promise, agreementId, agreementAmount]$$

In next steps Provider is exchanging promises with Accountant so in the end Accountant have promise to settle funds in Consumer's channel and Provider have payment promise to settle same amount of funds in Accountant's channel.

$$\Delta amount = newAgreementAmount - seenAgreementAmount$$

$$totalPromisedAmount = previousPromisedAmount + \Delta amount$$

Provider can do same promise exchange operations with many consumers at once. After any number of successful off-chain interactions, provider at any

point, in single transaction, can settle last payment promise with accumulated amount of payments done by couple of consumers (see figure 16. Additionally, if Provider is willing to take some risk (up to some amount) he can verify (exchange promises with accountant) not each payment and reduce amount of off-chain communication and fee paid to Accountant.

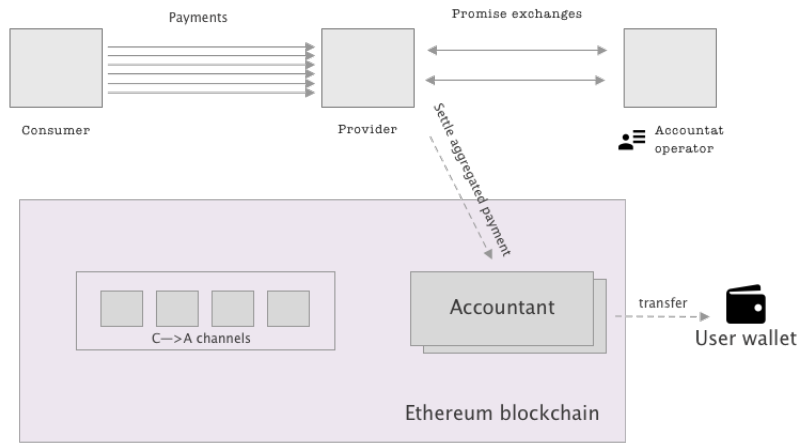


Figure 16: Provider settles aggregated promise

Accountant is also accumulating payment promises given by same Consumer to different Providers. When he will have payment promise with enough big value he can settle it on-chain (see figure 17) and rebalance paying and receiving channels.




Figure 17: Accountant rebalancing is simple settlement of collected promises

## 4.5 Incoming channel funds guaranties and trustless rebalance

All analysed micropayment channel based solutions have common problem of need to lock funds in channels with each party. In hub based solutions this problem is even bigger and requires so called *"rich hub"*. Thanks to possibility to take stake and users separation to consumers and providers it is possible to decrease this problem.

```
1     lockedFunds: uint256 # amount of accountant funds already
      locked in channels
2     struct Channel:
3         loan: uint256      # amount lendd by provider
4         balance: uint256   # amount available to settle
5         settled: uint256   # total amount settled by provider
6
7     @public
8     def rebalanceChannel(channelId: bytes32):
9         newBalance: uint256 = channels[channelId].loan
10        assert newBalance > channels[channelId].balance
11
12        channel: Channel = channels[channelId]
13
14        if(newBalance > channel.balance):
15            lockedFunds += channel.balance - newBalance
16            assert token.balanceOf(this) >= lockedFunds
17
18        channel.balance = newBalance
```

Provider's stake can be taken and lendd into Accountant and imidiately locked in his receiving payment channel. In this way, accountant will lock his own funds, which means that Provider can settle payment promises in value of up to his stake size.

Additional benefit is that Accountant can't actually withdraw and use these funds, so channel rebalance can be done by anyone, without accountant's signature.

"Always online" is big problem for payment channels. If Provider is offline or is under DDoS attack, he will be unable to post the finalized state (during dispute period) and Accountant can have possibility to take funds locked in channel.

Proposed method don't have such problem, because funds are double locked. Accountant can't leave channel with more funds than his own locked part of balance. So if Provider have unsettled promise for smaller amount than he lendd to Accountant, he is safe even if he will "dissapear" for really long period of time. He still be able to get back all funds promised to him.



## 4.6 Transaction maker

One of most not user friendly parts of transferring ERC20 tokens is that this operation requires to own ethers to pay for gas. Another problem in working with smart contracts on Ethereum is that in terms to call some function user have to add so called Payload (first 4 bytes of keccak of function's signature) which looks to him as gibberish and is not supported by many wallets and especially by exchanges.

In case of topups, this problem is solved thanks to construct described in 4.3 section. Unfortunately for identity registration and promise settlement, user's application (e.g. dVPN mobile app or provider's node) would need to always have ethers topuped there. This is very inconvenient and may cause many problems for both advanced users and beginners.




Figure 18: Communication with blockchain via Transaction maker

For identity registration and promise settlement we may take use of additional service which we call "*Transaction maker*". The role of this service is to get user's requests and send transactions into blockchain (see figure 18).

All smart contracts created for Accountant pattern are constructed in a way, so any their call can be made from any account simply by adding channel's operator signature as parameter for function call. Additionally we construct payment promises in a way so only one signature is needed and it is possible to settle them in any order with any gaps. This means that promises can be published openly without any risk of attack or losing funds.

There could be deployed as many such services as needed. Each network

participant could decide to use own "*Transaction maker*", or even there could be version of accountant operator software where such service is integral part.

Sending transactions into blockchain costs so users could cover that cost, but instead of using additional currency (ethers in case of Ethereum), they could pay in tokens.

There is potential misbehavior of "*Transaction maker*", when he is getting paid but didn't do his job, or misbehavior of users when they decide not pay after transaction was sent to blockchain by "*Transaction maker*". To resolve this problem, functions of smart contracts usually have fee parameter in them.

```
1  function settlePromise(uint256 _amount, bytes32 _lock,
2      uint256 _fee, bytes memory _signature) public {
3
4      require(_amount >= _fee);
5
6      ...
7
8      // Send tokens
9      token.transfer(
10         party.beneficiary,
11         _unpaidAmount.sub(_fee)
12     );
13
14     // Pay fee for transaction maker
15     if (_fee > 0) {
16         token.transfer(msg.sender, _fee);
17     }
18 }
```

In this way "*Transaction maker*" will get his reward only if he actually will send transaction and users have no chance not to pay.




Figure 19: Paying via Transaction maker flow

## 4.7 Properties of the protocol

This protocol was design to be integral part of some decentralised system and is not aimed to work as general purpose payment network used in many application and daily life, when each party in the same time can receive and send funds.

Main properties of protocol are:

- Easy paying channel topup (even from exchange).
- Payments aggregation  $\implies$  minimal on-chain tx amount.
- Fast settlement into blockchain without need for other party to be on-line and cooperate.
- Guaranty of incomming channel balance/collateral size.
- Simple off-chain communication (one roundtrip payments).
- Possibility to avoid accountant for at least part of transactions.
- Pull based interations with accountant  $\implies$  providers don't have to be externally exposed or maintain live connection with accountant.
- Secure  $\implies$  double spend protection using non-custodian accountant model.
- Fast channel opening  $\implies$  possibility not to wait until opening tx will be mined.
- Instant finality of payments.

## 5 Potential future work

There are three main directions in which future development of payment solution could evolve:

- migrate into Connexor or Perun;
- migrate into Plasma + channels;
- add support of counterfactual channels for accountant pattern.

## 5.1 Migration into Connex or Perun

Two most closed and most promising solutions to described Accountant pattern are *Perun virtual payment hubs* and *Connex hub*. While Perun team is still working on their first version of protocol implementation, Connex team already working on thirs iteration of their hub solution (this time using Counterfactual state channels). Unfortunately at the moment both solutions are not ready to use inside dVPN application written by Mysterium team.

In Connex case, they're changing their solution a lot, so it's worth to wait until their implementation will stabilise. Also their client application is written in TypeScript which means that Mysterium would need to reimplement it in goLang and take care of all protocol changes. Another problem is, that their off-chain messaging is much more complicated and requires storing a quite a lot of state by client nodes.

Perun team is creating their client node using goLang, which makes it good choice, but because protocol itself is complicated and implementation is in early stage, there is quite hard to say if it will work as expected and will provide all needed features for good user experience (described in 2.3 section, 7th requirement).

In the future, when these solutions will be more mature and if Accountat pattern will not serve as expected, there may be made migration into one of these solutions.

## 5.2 Plasma plus channels

In Ethereum community there are a lot of development going around various Plasma implementations. Unfortunately each analysed Plasma implementation had own issues (either centralisation, either mass exit risk, either having support only for non-fungible tokens). Additionally, as was described in 3.2 section, in initial state running own Plasma could be expensive, and there are no popular and decentralised Plasma implementation runned by thirs party teams.

In the future however, if there would be proposed Plasma implementation which solves earlier mentioned problems, would have already developed tooling which would help to run own implementaiton easily or they would appear decentralised and popular solution implemented and runned by third party team, Mysterium team could easily enough adapt their dVPN node software to use uni-directional, payment promise based channels on top of Plasma.

### **5.3 Take use of Counterfactual state channels**

Generalized state channels described in Counterfactual [2] white paper looks like very powerful solution. Also we see that there is some traction around this solution. If this trend will grow, in the future we may take some effort to research possibility to refactor Accountant pattern proposed uni-directional channels and implement them in counterfactual way. Ideally there could be possible to install them as an application on top of existing opened channel with counterfactual hub. This would help to have minimal changes inside dVPN node software but in same time be part of bigger network.

## References

- [1] Vitalik Buterin. *EIP1014: Skinny CREATE2*. <http://eips.ethereum.org/EIPS/eip-1014>.
- [2] Jeff Coleman, Liam Horne, and Li Xuanji. *Counterfactual: Generalized state channels*. 2018.
- [3] Stefan Dziembowski et al. “Perun: Virtual payment hubs over cryptocurrencies”. In: *2019 IEEE Symposium on Security and Privacy (SP)*. 2019, pp. 327–344.
- [4] Peter Murray, Nate Welch, and Joe Messerman. *EIP1167 Minimal Proxy Contract*. <http://eips.ethereum.org/EIPS/eip-1167>.
- [5] “Mysterium Network”. In: ().
- [6] *Mysterium payments smart contracts*. <https://github.com/mysteriumnetwork/payments-smart-contracts>. 2019.
- [7] J. Poon and T. Dryja. “The bitcoin lightning network: Scalable off-chain instant payments”. In: (2016).
- [8] Joseph Poon and Vitalik Buterin. “Plasma: Scalable autonomous smart contracts”. In: *White paper* (2017), pp. 1–47.